

Datathon

Hao Guo

February 2026

1 Model

1.1 Assumptions

- Single global half-life parameter for all user and all words (can be adjusted to improve power)
- independent factorization of difficulty via language, type of speech, and word length
- Approximate memory decay outside of Beta distribution

1.2 Derivation

The model we use to answer the question:

What is the probability \mathbb{P} that the user U still remembers the word W in language L , given W is in type T with length l after time t , where

$$T \in \{\text{noun, verb, number, adjective, pronoun, preposition, adverb}\},$$

and

$$L \in \{\text{English, Spanish, German, French, Italian, Portuguese}\}.$$

uses uses a Bayesian Beta–Binomial model and the posterior mean $\mathbb{P} = \frac{\alpha}{\alpha+\beta}$ as the recall probability.,

$$\mathbb{P} \rightarrow f(\mathbb{P} | x) = \frac{f(x | \mathbb{P})g(\mathbb{P})}{\int_{\mathbb{P}} f(x | \mathbb{P})g(\mathbb{P})d\mathbb{P}}$$

where $\mathbb{P} \sim \text{Beta}(\alpha, \beta)$.

Let α_p and β_p denote prior beliefs of α and β , then we can derive them as:

$$p_{prior} = \sigma(b_0 + b_L f(L) + b_T g(T) + b_l h(l))$$

where

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$

and

$$\begin{aligned}\alpha_p &= z \cdot p_{prior} \\ \beta_p &= z \cdot (1 - p_{prior})\end{aligned}$$

where f, g, h return difficulty related features (higher=easier), which are combined in a logistic regression to get the prior recall probability p_{prior}

$f(L)$ is computed based on studies FSI language difficulty,

$$f(L) = \epsilon + (1 - 2\epsilon) \left[1 - \frac{H(L) - H_{min}}{H_{max} - H_{min}} \right]$$

where ϵ is a tiny number that ensures $f(L) \in (0, 1)$, $H(L)$ returns the hours needed to learn language L .

$g(T)$ is retrieved from the study of Suranto [1], which describes the percentage of identified words in different word types for students.

$h(l)$ is estimated using a linear regression with rows with `history_seen ≤ 3`:

$$h(l) = \beta_0 + \beta_1 l$$

Once we have α_p and β_p , we obtain α and β by:

$$\begin{aligned}\alpha &= \text{History correct} + \alpha_p \\ \beta &= \text{History seen} - \text{History correct} + \beta_p\end{aligned}$$

and \mathbb{P} is calculated by:

$$\mathbb{P} = \frac{\alpha}{\alpha + \beta}$$

we can further apply a memory decay function

$$\begin{aligned}d(t) &= 2^{-t/h} \\ t &= \frac{\text{delta}}{3600} \text{ (unit in hours)}\end{aligned}$$

with the half life model. Then the probability \mathbb{P}_t of the user U remembering the word W after time t is calculated by:

$$\mathbb{P}_t = d(t) \cdot \mathbb{P}$$

2 Fitting

Once we have the model working, we need to fit our model, then test the correctness and evaluate. (How good is the model?)

2.1 Fit&Correctness

To fit our model, we use the following procedure:

1. randomly split users into train and test (say 80%:20%)
2. for each user, sort by timestamp (likely already done so)
3. use the train group to train the model, and test it with the test group

Then for each row i ,

- $n_i = \text{session_seen}$
- $m_i = \text{session_correct}$

compute Binomial log-likelihood:

$$\log Y_i = m_i \log \mathbb{P}_{ti} + (n_i - m_i) \log(1 - \mathbb{P}_{ti})$$

then calculate the average negative log-likelihood (NLL):

$$NLL = -\frac{1}{N} \sum_i \log Y_i$$

The smaller the NLL, the better the model describes the data.

Parameters we can tune to improve the model:

- z: adjust the weight of prior beliefs
- h: half life parameter
- $\vec{b} = \{b_0, b_L, b_T, b_l\}$

We want:

$$\theta^* = (z^*, h^*, \vec{b}^*) = \operatorname{argmin}_\theta NLL(\theta),$$

where $\theta = (z, h, \vec{b})$. We can use `scipy.optimize.minimize`.

2.2 Evaluation

We can compare our model to other common models. Namely:

- Half-Life Regression (HLR) (The ultimate model Duolingo uses)
- ACT-R Memory Model (A decent model)
- Logistic Regression / Neural Classifiers for Recall (We should better beat this)
- Simple Exponential Forgetting Curve (The baseline we must beat)

For testing, we can do a clustered sampling of our dataset (maybe 3 users per language) to speed up the evaluation.

We might not have time to implement these models, but we could mention them in the presentation.

3 Advantages and limitations

3.1 Limitations

- Likely not as accurate as the HLR model (obviously)
- Cannot fully utilize time data
- The linear regression used in $h(l)$ may not predict well for words with length exceeding the longest word in the dataset.

3.2 Advantages

- Computationally cheap
- Does not require a lot of data
- Since it is structurally simple, modifications can be made easily.

4 Test results

The result is obtained by running `model_stream.py` with parameters:

- `CHUNK_SIZE = 200_000` # The size of each chunk of data
- `RANDOM_SEED = 42`
- `EPS = 1e-9` # to avoid log(0)
- `EPS_F = 1e-3` # epsilon in $f(L)$
- `USE_SUBSAMPLE_FOR_TRAINING = True`
- `TRAIN_SUBSAMPLE_ROWS = 1_000_000` # rows used per NLL eval while fitting
- `MAXITER = 50` # optimizer iterations
- `PRIOR_SAMPLE_MAX = 1_000_000` # sample size for estimating the prior
- half-life upper bound = 20000.0

Note that for performance improvements, \vec{b} is optimized first, then (z, h) follows.

```
Pass 1: collecting unique users...
Total unique users: 115222
Train users: 92177, Test users: 23045
Pass 2: estimating h(l) = beta0 + beta1*l via streaming regression...
h(l): beta0 = 0.8880, beta1 = 0.0013 (from 3415409 rows)
Pass 3: sampling low-history rows to fit prior coefficients (b0, bL, bT, b1)...
```

```
Prior logistic regression sample size (before cleaning): 1046364
Prior logistic regression sample size (after cleaning): 1045821
Prior coefficients (from GLM):
b0 = 11.4914, bL = 0.3322, bT = 0.0086, bl = -10.1882
Pass 4: precomputing numeric feature file (this is a one-time cost)...
Feature file written to: precomputed_features.csv
Fitting z and half-life h on training users (using subsample for speed)...
```

Optimization result for z and h:

```
message: CONVERGENCE: NORM OF PROJECTED GRADIENT <= PGTOL
success: True
status: 0
fun: 0.5604880087063433
x: [ 2.743e+01  1.560e+04]
nit: 26
jac: [ 1.588e-06 -6.661e-08]
nfev: 90
njev: 30
hess_inv: <2x2 LbfgsInvHessProduct with dtype=float64>
```

Estimated z = 27.4349

Estimated h = 15604.8663 hours (half-life)

Computing final NLL on full train and test sets (this may take a bit)...

```
Final NLL (train): 0.565432
Final NLL (test) : 0.569647
```

h(l): beta0 = 0.8880, beta1 = 0.0013

$\beta_0 = 0.888$: baseline easiness at length 0

$\beta_1 = 0.0013$: easiness increases by this number per additional character

Users are slightly more likely to get longer words correct.

b0 = 11.4914, bL = 0.3322, bT = 0.0086, bl = -10.1882

Different languages and word type barely have any effect. Word length seem to have huge effect, but it is mostly canceled by b0.

Estimated z = 27.4349

The prior behaves like 27 "unseen" trials.

h = 15604.8663 hours

Memory basically does not decay.

Final NLL (train): 0.565432

Final NLL (test) : 0.569647

tiny gap between train and test meaning the model is not overfitting.

NLL (test) implies a strong model with correctness = 0.743 (if we guess randomly at each trial (correctness = 0.5), we would get NLL = 0.693)

Overall the coefficients seem a bit weird, but the model works fine. Probably due to Duolingo's learning method, i.e. user can tap on the word and see translation.

As a comparison, HLR results in NLL = 1.052842, which is also weird since the model should not have behaved this bad.

All these cues may imply bugs in our code.

References

- [1] Suranto and Yuspik. "The Analysis of Student's Ability to Identify Parts of Speech". In: *English Teaching and Applied Linguistics Journal* 1.1 (2024), pp. 18–26.